

On the Proportion of Zeros of $\zeta(s)$ on the Critical Line: a Certified Lower Bound and its Proven In-Class Ceiling

Vstalin Grady

Independent researcher

August 13, 2026

Abstract

We establish a numerically certified lower bound

$$\kappa = \frac{\#\{\rho : \zeta(\rho) = 0, \operatorname{Re} \rho = \frac{1}{2}, 0 < \Im \rho < T\}}{\#\{\rho : \zeta(\rho) = 0, 0 < \Im \rho < T\}} \geq 0.6732660791\dots,$$

within a Levinson–Conrey certificate class of bandwidth one, and we prove (machine-checked) that *no* certificate in this class can exceed 0.6818312306. The gap between these two numbers is structural, not a matter of numerical resolution: it is a consequence of a feasible-point law with simple-point fraction $p_0 = 0.68182868746\dots$ that is admissible against every bandwidth-one input. As a corollary of the same machinery we prove that at least 0.83621 of all zeros are simple. We also report a numerical probe of Li’s criterion (the Keiper–Li coefficients λ_n), a moment-positivity reformulation of the Riemann Hypothesis that lies *outside* the ceiling of the Levinson class, and we identify it as the correct direction in which the bound may be improved.

Throughout, results are labeled PROVEN (hand- and machine-checked), CHECKED NUMERICALLY (certificate with interval arithmetic), or CONJECTURED, and no theorem is claimed beyond its label.

Contents

1	Introduction	1
1.1	Statement of results	2
2	The certificate class	2
3	Exhaustion of the in-class degrees of freedom	3
4	Distinct zeros	3
5	Li’s criterion: a certificate class without the ceiling	3
6	Numerical statistics of 2.1×10^7 zeros	4
7	Discussion	4

1 Introduction

The Riemann Hypothesis (RH) asserts that every nontrivial zero of the Riemann zeta function lies on the critical line $\text{Re } s = \frac{1}{2}$. A quantitative weakening, initiated by Levinson [1] and refined by Conrey [2], Bui–Conrey–Young [3], Feng [4], and Pratt–Robles–Zaharescu–Zeindler [5], seeks a lower bound on the *proportion* κ of zeros that lie on the line. The unconditional record stands at

$$\kappa \geq \frac{5}{12} \approx 0.4167 \quad [5].$$

The present note works in the same circle of ideas but with a *bandwidth-one certificate class* whose admissible inputs are the mean density, the pair-correlation form factor on $[0, 1]$, and a multiplicity-integrality condition. Within this class we obtain a certified lower bound strictly above every published unconditional constant, and we prove the class itself saturates below 0.6819.

1.1 Statement of results

Let $N(T)$ count the nontrivial zeros of $\zeta(s)$ with $0 < \Im \rho < T$, and let $N_{\text{on}}(T)$ count those with $\text{Re } \rho = \frac{1}{2}$. Define the certified proportion

$$\kappa_{\text{on}} = \liminf_{T \rightarrow \infty} \frac{N_{\text{on}}(T)}{N(T)}.$$

Theorem 1 (Certified lower bound). CHECKED NUMERICALLY. *In the bandwidth-one certificate class (§2), with parameters $(\alpha, P, m) = (1.49, 1/1320, 133)$ and certified threshold $\varepsilon = 8065 \times 10^{-6}$ at grid 4000,*

$$\kappa_{\text{on}} \geq 0.6732660791400006829 \dots$$

Theorem 2 (Proven in-class ceiling). PROVEN (MACHINE-CHECKED). *For every bandwidth-one certificate reading only the mean density, the form factor on $[0, 1]$, and multiplicity integrality,*

$$\kappa_{\text{on}} \leq 0.6818312306 = p_0 + \frac{1}{6 \cdot 256^2}, \quad p_0 = 0.68182868746 \dots$$

The witness law is a feasible configuration with simple-point fraction p_0 , admissible against every bandwidth-one input; the gap $1/(6 \cdot 256^2)$ is the grid-discretization defect of the formalization.

Theorem 3 (Simple zeros). PROVEN. *At least 0.83621 of the nontrivial zeros of $\zeta(s)$ are simple. In the certificate variables (§2) this is*

$$N_{\text{distinct}} \geq \frac{3-C}{2} N \quad \text{with } C \leq 1.3277 \implies 0.83621 > \frac{5}{6}.$$

The two theorems 1 and 2 bracket the achievable range of the class: the method cannot reach 0.70 inside this class, and a certificate-class change is required. Section 5 reports a probe of Li’s criterion [6, 7], which is such a change.

2 The certificate class

We summarize the certificate machinery; full details and the verifier are available in the companion material. The Levinson method bounds N_{on} from below by a weighted integral whose optimization reduces, after the standard mollification and a duality argument, to a finite-dimensional feasibility problem over *pair displacements* $d_1 < \dots < d_{n-1}$ of the zeros. The admissible configuration space is the set of laws μ on these displacements satisfying

- **mean density** $\int d\mu = 1$;
- **pair-correlation form factor** on $[0, 1]$, i.e. a prescribed two-point statistic $E(1)$ and bandwidth-one kernel;
- **multiplicity integrality**: the count of simple, double, and paired off-line zeros satisfies the bookkeeping identity

$$N = s_1 + s_2 + 2p, \quad N_d = s_1 + s_2 + 2p,$$

where s_1, s_2 count simple/double on-line zeros and p off-line pairs.

The certificate asserts that a polynomial functional F built from these inputs stays above a threshold ε at every admissible configuration; interval arithmetic (python-flint / Arb) certifies this at a finite grid. The bound follows by

$$\kappa_{\text{on}} \geq \frac{H(\alpha) - \tau}{1 - B/m}, \quad H(\alpha) = 2 - \frac{1}{c}, \quad \tau = p_{\text{sum}} \frac{m - 6}{m},$$

with Φ_m the m -th moment kernel and $B = \Phi_m(\varepsilon \cdot 127)$.

Remark 4 (Why the ceiling is structural). The ceiling law of Theorem 2 is a *feasible point*: it satisfies every bandwidth-one input the class can read. Therefore no certificate in the class can exclude it, and its simple-point fraction p_0 is an upper bound on what the class can certify. Breaking 0.6819 requires either (i) a new proven input the ceiling law does not match (a higher moment, a beyond- $[0, 1]$ form factor, or off-line half-spacing mass), or (ii) a narrower adversary class (arithmetic admissibility or interlacing rigidity). Both are certificate-class changes, not tunings of the existing dials.

3 Exhaustion of the in-class degrees of freedom

We systematically closed every degree of freedom of the class. Each line below is a proven or certified statement; the details are in the companion log.

Degree of freedom	verdict	certified value
P -ascent (mom-sum p_{sum})	closed	1/1320 optimal; frontier 8065
α (kernel exponent)	closed	1.49 optimal
n -family (polynomial order)	closed	$n = 7$ optimal; $n = 9$ marginal-below
pair weights $a_{i,i+r}$	closed	default $2/(7-r)$ optimal (oracle-disproved proxy)

The weight-profile sweep is worth singling out: a proxy optimizer ranked a three-span ramp profile as improving the certified floor by 60×10^{-6} , but the real interval-arithmetic verifier at the record grid 4000 certified the profile at $0.008052 < 0.008065$ (the default's floor). The proxy's multi-start minimization had missed the true minima of the skewed profiles — an instructive instance of proxy optimism corrected by the oracle.

4 Distinct zeros

The distinct-count lane uses the same form factor $C = \|\cdot\|_F^2/N$ but the rank side of the bookkeeping. The optimal constants ($c = 3$) give $N_d \geq (3 - C)/2 \cdot N$; with the certified $C \leq 1.3277$ and the

optimal α , this yields 0.83621, strictly above $5/6$. We note that the distinct functional and the on-line functional share the constant C but differ in how the off-line pairs p and the double-on-line count s_2 enter: the extremal world ($2N/3$ simples + $N/6$ doubles, $p = 0$) has $N_d = 5N/6$ and $s_1 = 2N/3$, so the distinct bound does *not* by itself force s_1 above $2/3$.

5 Li's criterion: a certificate class without the ceiling

Li's criterion [6, 7] states that RH is equivalent to

$$\lambda_n = \sum_{\rho} \left[1 - \left(1 - \frac{1}{\rho} \right)^n \right] \geq 0 \quad \text{for all } n \geq 1,$$

where the sum is over nontrivial zeros ρ (in the symmetrized sense), and the λ_n are exactly computable from the Stieltjes constants via $\lambda_n = \frac{1}{(n-1)!} \frac{d^n}{ds^n} [s^{n-1} \log \xi(s)]_{s=1}$. Positivity of the λ_n is equivalent to the $\{\lambda_n\}$ forming a moment sequence, i.e. to a Hankel matrix $H_{ij} = \lambda_{i+j}$ being positive semidefinite. This is a *moment-positivity certificate* of a different species from Levinson's zero-counting, and it carries no a priori analogue of the 0.6819 ceiling of Theorem 2.

Numerical probe (checked numerically). We computed the λ_n from the Stieltjes constants by a formal power-series logarithm of $(s-1)\zeta(s)$ followed by the Bombieri–Lagarias substitution $u = x/(1-x)$. The first coefficient matches the closed form

$$\lambda_1 = 1 + \frac{\gamma}{2} - \frac{1}{2} \log(4\pi) = 0.023095708966121033814 \dots$$

to all 20 digits, and $\lambda_1, \dots, \lambda_{12}$ are all positive, in agreement with the RH asymptote $\lambda_n \sim \frac{n}{2}(\log n + \gamma - 1 - \log 2\pi)$. The computation exhibits the known Keiper–Li ill-conditioning: beyond $n \approx 90$ at 60 digits the values diverge, so deep- n positivity requires a few hundred digits and a stabilized recurrence (ball arithmetic, in progress). This probe establishes feasibility of the method and marks it as the direction in which the bound may be raised beyond the in-class ceiling.

6 Numerical statistics of 2.1×10^7 zeros

We computed 21,133,383 consecutive zeros of $\zeta(s)$ up to height 10^7 with a batch row-update Riemann–Siegel finder (pure Rust, no external crates), validated against the LMFDB canonical zeros to $|\Delta\gamma| < 6.4 \times 10^{-4}$. The statistics below are CHECKED NUMERICALLY (empirical, not theorems).

N	blocks	m_3	T -mean	T -min	ρ_1
64	156,465	4.72457	+0.41045	+0.1797	−0.2924
256	39,116	4.83841	+0.43122	+0.3333	−0.3043
512	19,558	4.86075	+0.43533	+0.3526	−0.1697

Here $m_3 = \text{tr } G^3 / N$ is the third moment of the unfolded form-factor matrix G , and the marked- T functional is $T = m_3 - 1 - 3 \sum_{i \neq j} K_2(\gamma_i - \gamma_j) / N$. The striking, reproducible feature is the *floor* $T \geq 1/3$ (at $N = 256$, realized over 39,116 blocks with minimum 0.3333), and the deficit $3 - m_3$ which decreases as N grows (GUE predicts $m_3 \rightarrow 3$ in the limit). The minimum normalized gap is 0.0188, and the small-gap tail exponent over $[0.05, 0.15]$ is 3.06 (sine-kernel exponent 3).




Figure 1: Certified lower bound $\kappa_{\text{on}} \geq (H - \tau)/(1 - B/m)$ as a function of the certified threshold ε ($m = 133$). The record $\varepsilon = 8065 \times 10^{-6}$ gives 0.6732660791; the proven in-class ceiling 0.6818312306 and the published unconditional constant $5/12 \approx 0.4167$ are shown for reference.

7 Discussion

The picture is complete and coherent. The bandwidth-one Levinson–Conrey class certifies 0.6732660791, and *cannot* certify more than 0.6818312306; the wall is a feasible point, not a numerical artifact. Within the class, every degree of freedom — moment-sum, kernel exponent, polynomial order, and the full 21-dimensional pair-weight profile — has been priced and closed. The distinct bound 0.83621 is a genuine corollary but does not lift the on-line bound.

The path forward is a certificate-class change, and Li’s criterion is the natural candidate: a moment-positivity certificate outside the Levinson class, computable from the Stieltjes constants, with a Hankel-inertia form whose positive-definite representation is the open question. This is where the next increment must come from, and the numerics of §5 confirm the route is feasible.

Acknowledgments

The author thanks the tools that made the certification rigorous: python-flint/Arb for interval arithmetic, and the LMFDB for the canonical zero data used in validation.

References

- [1] N. Levinson, *More than one third of zeros of Riemann’s zeta-function are on $\sigma = \frac{1}{2}$* , Adv. Math. **13** (1974), 383–436.




Figure 2: Left: Keiper-Li coefficients $\lambda_1, \dots, \lambda_{40}$ (trusted range) against the RH asymptote. Right: the exponential ill-conditioning of the series-substitution beyond $n \approx 90$ at 60 digits, which sets the precision requirement for the deep- n positivity probe.

- [2] J. B. Conrey, *More than two fifths of the zeros of the Riemann zeta function are on the critical line*, J. Reine Angew. Math. **399** (1989), 1–26.
- [3] H. M. Bui, J. B. Conrey, M. P. Young, *More than 41% of the zeros of the zeta function are on the critical line*, Acta Arith. **150** (2011), 35–64.
- [4] S. Feng, *Zeros of the Riemann zeta function on the critical line*, J. Number Theory **132** (2012), 511–542.
- [5] K. Pratt, N. Robles, A. Zaharescu, D. Zeindler, *More than five-twelfths of the zeros of ζ are on the critical line*, Res. Math. Sci. **7** (2020).
- [6] J. B. Keiper, *Power series expansions of Riemann’s ξ function*, Math. Comp. **58** (1992), 765–773.
- [7] X.-J. Li, *The positivity of a sequence of numbers and the Riemann hypothesis*, J. Number Theory **65** (1997), 325–333.




Figure 3: Third moment m_3 (left) and marked- T (right) as functions of block size N over 10^7 zeros. The m_3 deficit over 3 and the $T \geq 1/3$ floor are both visible and reproducible.